GLOBAL ENTERPRISE POLICY: AI GOVERNANCE AND HUMAN INTERVENTION

Document ID: AI-POL-2026-09

Effective Date: March 1, 2026

Last Reviewed: July 15, 2026

Owner: Office of the Chief Data Officer

SECTION 1: HUMAN-IN-THE-LOOP MANDATE

Autonomous AI agents are strictly forbidden from executing or assigning task items automatically to production database records. All extracted action items, risk assessments, and schedule modifications generated by LLM pipelines require explicit manual review, modification, and authorization from a human operator prior to final persistence. Any system that bypasses this human-in-the-loop requirement shall be treated as a Severity 1 compliance violation.

SECTION 2: MODEL TRANSPARENCY AND AUDIT

All AI models used in decision-support pipelines must have their model name, version, provider, and configuration parameters documented in a system-level audit record. This includes the model architecture (e.g., Qwen, LLaMA, Mistral), parameter count, quantization level, and context window size. Audit records must be retained for a minimum of 24 months. At least once per quarter, a designated compliance officer must review the audit trail for anomalies. The system must produce an automated quarterly AI usage report listing all models invoked, total inference calls, and any schema validation failures or unparseable outputs encountered.

SECTION 3: STRUCTURED OUTPUT AND DATA QUALITY

LLM pipelines that produce structured JSON output for downstream consumption must enforce strict schema validation at the application layer. Raw, unvalidated LLM output must never be passed directly to frontend rendering components. Every extraction pipeline must implement a Pydantic or equivalent schema binding, and any output that fails validation must be logged as an error event and surfaced to a human operator for remediation — never silently discarded. The schema must include mandatory source-attribution fields linking every extracted item to its originating document and section.

SECTION 4: FAIRNESS AND BIAS REVIEW

Before deploying any LLM-based extraction pipeline to production, the project team must conduct a fairness review assessing whether the model exhibits bias in classifying content from different document types, authorship styles, or domain vocabularies. The review must include at least 20 sample documents spanning multiple authors, departments, and writing styles. Results must be documented and signed off by the project lead.

SECTION 5: PROHIBITED USE CASES

AI extraction pipelines shall not be used for: (a) automated hiring or promotion decisions, (b) automated generation of legally binding contracts, (c) automated medical diagnosis or treatment recommendations, (d) any use case where an incorrect extraction could cause financial loss exceeding $10,000 without human review. Violations of this section are grounds for immediate project suspension.

SECTION 6: COMPLIANCE CHECKLIST FOR AI PROJECTS

Every AI project MUST complete the following before GA launch:

- AI-REG-01: Register the model name, version, and provider in the corporate model registry.
- AI-REG-02: Document the human review workflow — screenshots or flowcharts showing where human intervention occurs.
- AI-REG-03: Implement structured output validation with error logging.
- AI-REG-04: Complete fairness and bias review with documented results.
- AI-REG-05: Schedule quarterly audit reviews for the first year of operation.
- AI-REG-06: Define an escalation path for when the LLM produces unparseable output.
- AI-REG-07: Publish a user-facing disclaimer explaining that AI-extracted content requires human verification.