GLOBAL ENTERPRISE POLICY: DATA RETENTION AND VECTOR STORAGE

Document ID: DAT-POL-2026-04
Effective Date: February 15, 2026
Last Reviewed: July 22, 2026
Owner: Data Governance Office

SECTION 1: STORAGE CONTAINMENT

All structured extraction summaries and user metadata must reside within approved database systems. Vector embeddings and semantic document indices must be contained within pgvector instances operating inside the primary PostgreSQL infrastructure. External third-party cloud vector stores (e.g., Pinecone, Weaviate Cloud, Qdrant Cloud) are strictly prohibited. Any proposal to use an external vector database requires a formal data residency review and written approval from the Data Governance Office.

SECTION 2: TEMPORARY FILE PURGING

Uploaded raw documents (including PDF, DOCX, and TXT files) processed during multi-file document ingestion pipelines must be securely purged from temporary application server storage within 24 hours of successful vectorization. The purge must use secure deletion methods (file overwrite before unlink, or equivalent). Logs confirming successful purge operations must be retained for audit purposes. Retaining raw uploaded files beyond 24 hours without an explicit data retention waiver is a compliance violation.

SECTION 3: VECTOR DATA RETENTION

Document chunk embeddings stored in pgvector must be retained for the duration of the associated project's lifecycle plus 90 days. After project closure, all associated chunks, embeddings, and metadata must be permanently deleted. Projects that remain inactive (no new uploads or analyses) for more than 180 days shall be flagged for archival review. The system must provide a project archival API endpoint that purges all associated data in a single atomic operation.

SECTION 4: USER DATA AND EXTRACTION RESULTS

Structured extraction results (facts, decisions, risks, action items, etc.) that have been reviewed and approved by a human operator must be retained for a minimum of 3 years for audit and compliance purposes. Unreviewed or draft extraction results that have not been approved within 30 days of generation may be automatically purged. Users must be notified via email 7 days before unreviewed results are purged.

SECTION 5: ENCRYPTION REQUIREMENTS

All data at rest within PostgreSQL must be encrypted using AES-256 or equivalent. Data in transit between services (frontend to backend, backend to database, backend to Ollama) must use TLS 1.3 or higher. Vector embeddings stored in pgvector are considered sensitive data and must be encrypted at rest with the same standard as user PII.

SECTION 6: CROSS-BORDER DATA TRANSFER

No project data — including document text, embeddings, extraction results, or user metadata — may be transferred to servers or services located outside the organization's designated data residency region without explicit approval from the Data Governance Office. All third-party API integrations must be reviewed for data residency compliance before integration. This includes any LLM provider APIs that may process data on external servers.

SECTION 7: RETENTION COMPLIANCE CHECKLIST

Every data-handling application must complete the following before GA:

- DAT-REG-01: Document the data retention schedule for raw uploads, embeddings, and extraction results.
- DAT-REG-02: Implement automated temporary file purging within 24 hours of vectorization.
- DAT-REG-03: Implement project archival with atomic data deletion.
- DAT-REG-04: Verify AES-256 encryption at rest for all PostgreSQL data.
- DAT-REG-05: Verify TLS 1.3 for all inter-service communication.
- DAT-REG-06: Complete data residency review confirming no cross-border transfers.
- DAT-REG-07: Implement user notification for pending data purges (7-day warning).
- DAT-REG-08: Schedule quarterly data retention audits for the first year.